

Final Project Report Template

1. Introduction

1.1. Project overviews

The project titled "Visual Diagnostics: Detecting Tomato Plant Diseases Through Leaf Image Analysis" aims to develop an accurate and efficient deep learning model for diagnosing various tomato plant diseases from leaf images. The primary objective is to assist farmers in the early detection and management of plant diseases, thereby improving crop yield, reducing economic losses, and promoting sustainable farming practices. Given the susceptibility of tomato plants to numerous diseases, early and precise diagnosis is crucial for effective management. Traditional diagnostic methods rely heavily on expert knowledge and laboratory tests, which are often time-consuming and costly. This project leverages advanced machine learning techniques to provide a cost-effective and accessible diagnostic tool for farmers. Utilizing a dataset sourced from Kaggle, which includes images of tomato leaves categorized into ten types—ranging from healthy leaves to those affected by diseases such as spider mites and yellow leaf curl—the project aims to harness the power of image analysis for sustainable agricultural practices.

1.2. Objectives

The primary objective of this project is to develop a robust and efficient deep learning model capable of accurately diagnosing various tomato plant diseases through the analysis of leaf images. This model aims to provide farmers with a powerful diagnostic tool that facilitates early detection and effective management of plant diseases. By leveraging advanced machine learning techniques, the project seeks to overcome the limitations of traditional diagnostic methods, which are often dependent on expert knowledge and costly, time-consuming laboratory tests. The goal is to create an accessible, cost-effective solution that improves crop health, enhances yield, and reduces economic losses for farmers. Ultimately, the project strives to promote sustainable farming practices by equipping farmers with a reliable and easy-to-use tool for disease management.

2. Project Initialization and Planning Phase

2.1. Define Problem Statement

The project addresses the challenge of accurately diagnosing tomato plant diseases, which traditionally relies on expert knowledge and costly laboratory tests, by developing a deep learning model that can analyze leaf images and provide quick, reliable disease identification to support early and effective disease management for farmers.

2.2. Project Proposal (Proposed Solution)

The proposed solution involves developing an intuitive diagnostic tool that leverages advanced deep learning techniques to analyze images of tomato leaves and accurately identify various plant diseases. Using a pre-trained ResNet152V2 model fine-tuned with custom layers, the tool will process and classify leaf images into ten distinct categories, including healthy and diseased conditions. This model will be integrated into a user-friendly application, accessible via smartphones or dedicated devices, providing farmers with instant results and actionable treatment recommendations. By offering a reliable, cost-effective, and efficient method for disease management, this solution aims to reduce reliance on subjective visual inspections and expensive lab tests, thereby enhancing crop health, increasing yields, and improving farmer satisfaction.

2.3. Initial Project Planning

The initial project planning phase involves a comprehensive approach to set the foundation for developing the diagnostic tool. It begins with collecting and organizing a diverse dataset of tomato leaf images, categorized by different diseases, sourced from platforms like Kaggle. Following this, the deep learning model architecture is defined, utilizing the pre-trained ResNet152V2 model with custom layers tailored for disease classification. The project plan includes setting up a robust training environment on Kaggle to leverage GPU and TPU accelerators for efficient model training. Essential steps such as data preprocessing, augmentation, and normalization are incorporated to enhance model performance. The development process is structured into multiple sprints, with clear milestones and deliverables, ensuring systematic progress. Collaboration among team members, including data scientists, developers, and domain experts, is emphasized to ensure the project's success. Regular evaluations and iterations are planned to fine-tune the model and address any emerging issues, ultimately leading to the deployment of an accessible and reliable diagnostic tool for farmers.

3. Data Collection and Preprocessing Phase

3.1. Data Collection Plan and Raw Data Sources Identified

The data collection plan focuses on sourcing a comprehensive dataset of tomato leaf images, capturing a wide range of disease conditions to ensure robust model training. Images are collected from Kaggle's Tomato Leaf Dataset, which includes categories such as healthy, spider mites, yellow leaf curl, and others, each organized into respective subdirectories. The dataset, available in image format, is accessible publicly, providing a valuable resource for developing the diagnostic tool. Additionally, to enhance model accuracy, plans include gathering more images to address class imbalances and improve representation of all disease types. This extensive collection process is crucial for training a deep learning model capable of accurately diagnosing diseases from varied visual symptoms. The project will utilize Kaggle's platform for data storage and model training, leveraging its computational resources like GPUs and TPUs to handle the large dataset efficiently. This structured data collection approach ensures the model is trained on high-quality, diverse data, laying the groundwork for effective disease diagnosis.

3.2. Data Quality Report

The Data Quality Report highlights the assessment of the collected tomato leaf dataset to ensure its suitability for training the deep learning model. The primary issue identified is the imbalance in class distribution, with certain disease categories having significantly

fewer images, which could affect the model's performance. To address this, a data augmentation strategy is planned to generate synthetic samples for underrepresented classes, thereby balancing the dataset. Additionally, image quality was evaluated, and any low-resolution or noisy images were flagged for removal or enhancement to maintain data integrity. The dataset's organization into clearly labeled subdirectories based on disease type was verified, ensuring that data is correctly categorized for model training. Overall, the data quality assessment emphasizes resolving class imbalances and ensuring high-quality images, which are crucial for developing an accurate and reliable diagnostic tool. These measures are detailed in the resolution plan, ensuring systematic identification and rectification of data discrepancies to support the project's objectives.

3.3. Data Preprocessing

Data preprocessing is a critical step in preparing the tomato leaf dataset for training the deep learning model. This process begins with organizing the images into well-defined categories based on disease types, ensuring a clear structure for model input. Each image is then resized to a uniform dimension of 256x256 pixels to maintain consistency and optimize computational efficiency. Rescaling is performed by normalizing pixel values to a range of 0 to 1, which aids in faster convergence during training. Data augmentation techniques, such as rotation, flipping, and zooming, are applied to artificially expand the dataset and mitigate class imbalances, enhancing the model's ability to generalize across different disease conditions. Furthermore, the dataset is split into training, validation, and test sets to evaluate the model's performance accurately. Any low-quality or corrupted images identified during the data quality assessment are excluded to maintain high data integrity. This comprehensive preprocessing pipeline ensures that the input data is clean, consistent, and representative, providing a solid foundation for training an effective disease diagnostic model.

4. Model Development Phase

4.1. Model Selection Report

The Model Selection Report outlines the process of choosing the most suitable deep learning architecture for the tomato plant disease diagnostic tool. After evaluating several candidate models, including VGG16, InceptionV3, and ResNet152V2, the ResNet152V2 model was selected due to its superior performance in image classification tasks and its ability to learn intricate patterns through its deeper architecture. The model was pre-trained on the ImageNet dataset, providing a strong starting point with robust feature extraction capabilities. Custom layers were added to fine-tune the model specifically for tomato leaf disease classification, including a GlobalAveragePooling2D layer to reduce spatial dimensions, followed by dense layers for learning disease-specific features, and a final softmax layer for multi-class classification. The first 140 layers of ResNet152V2 were frozen to retain pre-learned features, while the remaining layers were made trainable to adapt to the new dataset. This combination of pre-trained and custom layers ensures a balanced approach, leveraging existing knowledge while tailoring the model to the specific problem. The Adam optimizer and categorical crossentropy loss function were selected to optimize training, focusing on accuracy as the primary metric. This model selection process, based on rigorous evaluation and customization, positions ResNet152V2 as the optimal architecture for achieving accurate and reliable tomato plant disease diagnosis.

4.2. Initial Model Training Code, Model Validation and Evaluation Report

The initial model training began with loading the ResNet152V2 architecture pre-trained on ImageNet and customizing it for tomato plant disease classification. The model was compiled with the Adam optimizer and categorical crossentropy loss function to optimize for accuracy. Training utilized a dataset split into training and validation sets, with data augmentation techniques applied to enhance model robustness. During training, performance metrics such as accuracy and loss were monitored to assess model convergence and generalization capability. After training, the model was evaluated on a separate test set to measure its effectiveness in accurately diagnosing various diseases from tomato leaf images. The validation and evaluation processes aimed to validate model performance against ground truth labels, ensuring its readiness for deployment as a reliable diagnostic tool for farmers.

5. Model Optimization and Tuning Phase

5.1. Tuning Documentation

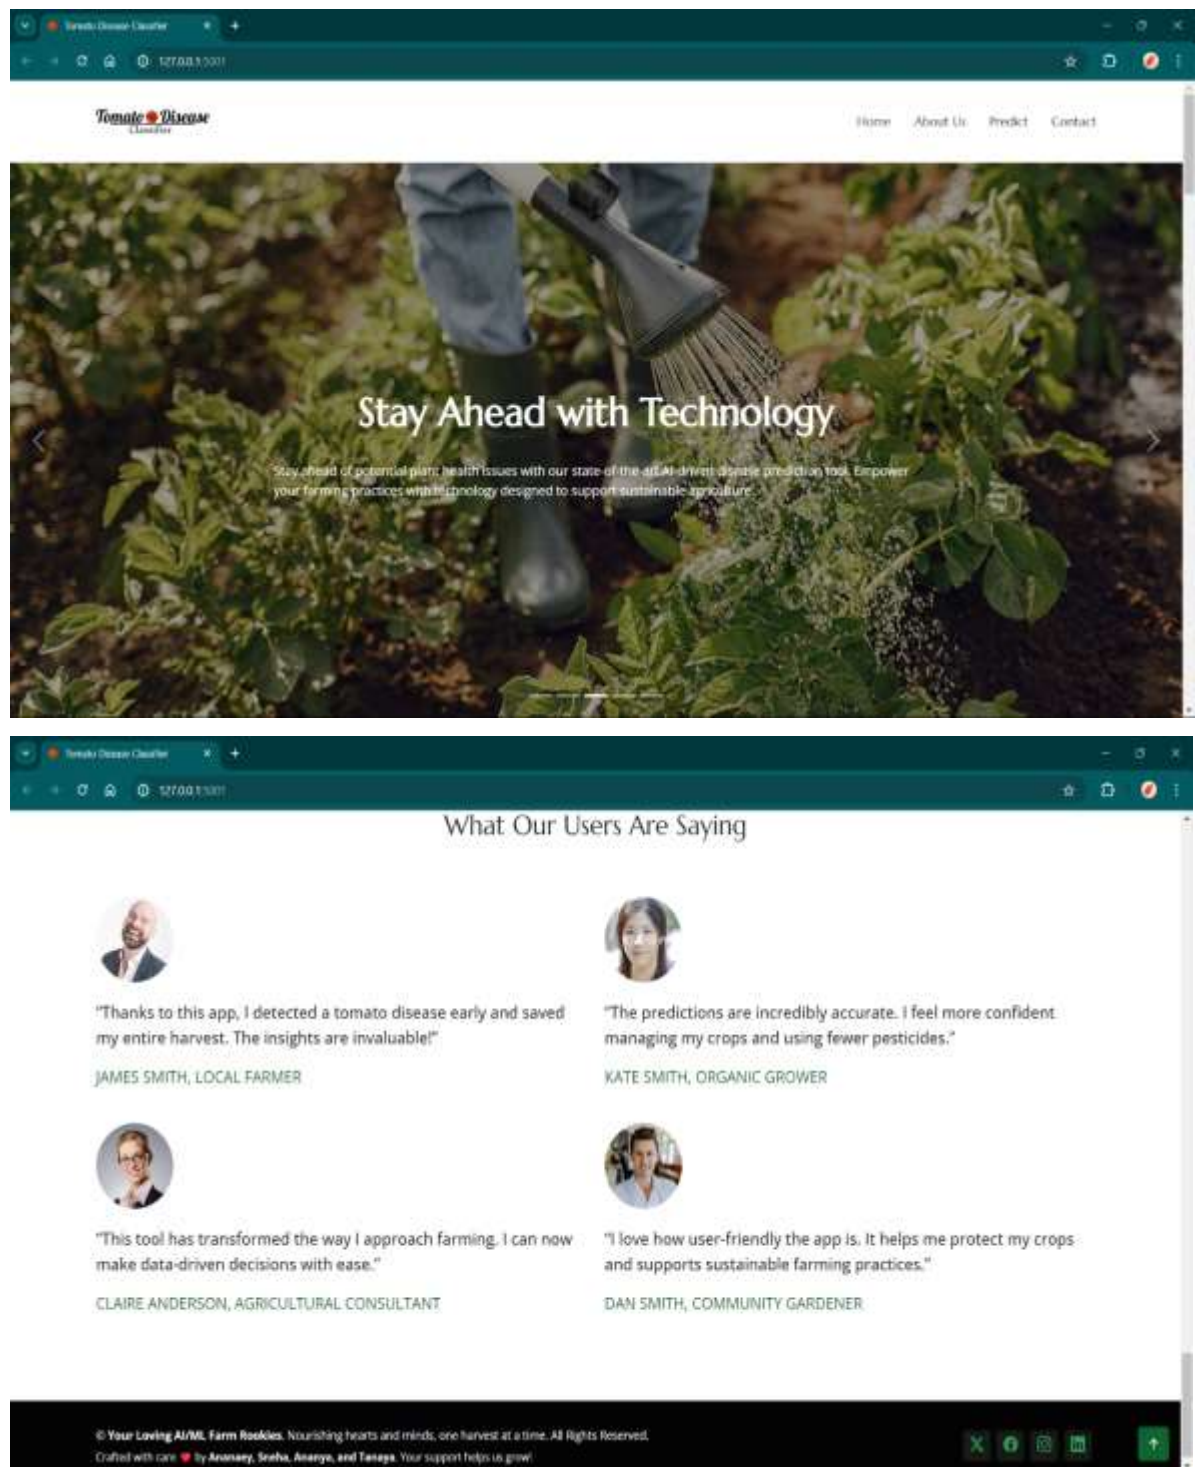
The Tuning Documentation details the process of optimizing the ResNet152V2-based model for enhanced performance in tomato plant disease classification. It begins with fine-tuning hyperparameters such as learning rate, batch size, and number of epochs through iterative experimentation. Techniques like learning rate schedules and early stopping criteria were implemented to prevent overfitting and achieve optimal convergence. Additionally, the impact of different data augmentation strategies was explored to improve model generalization across diverse disease conditions. Regular validation checks were conducted to monitor metrics such as accuracy and loss, ensuring the model's robustness and reliability in real-world scenarios. The tuning documentation captures these iterative adjustments and their outcomes, providing insights into the strategies that led to the final optimized configuration of the diagnostic tool.

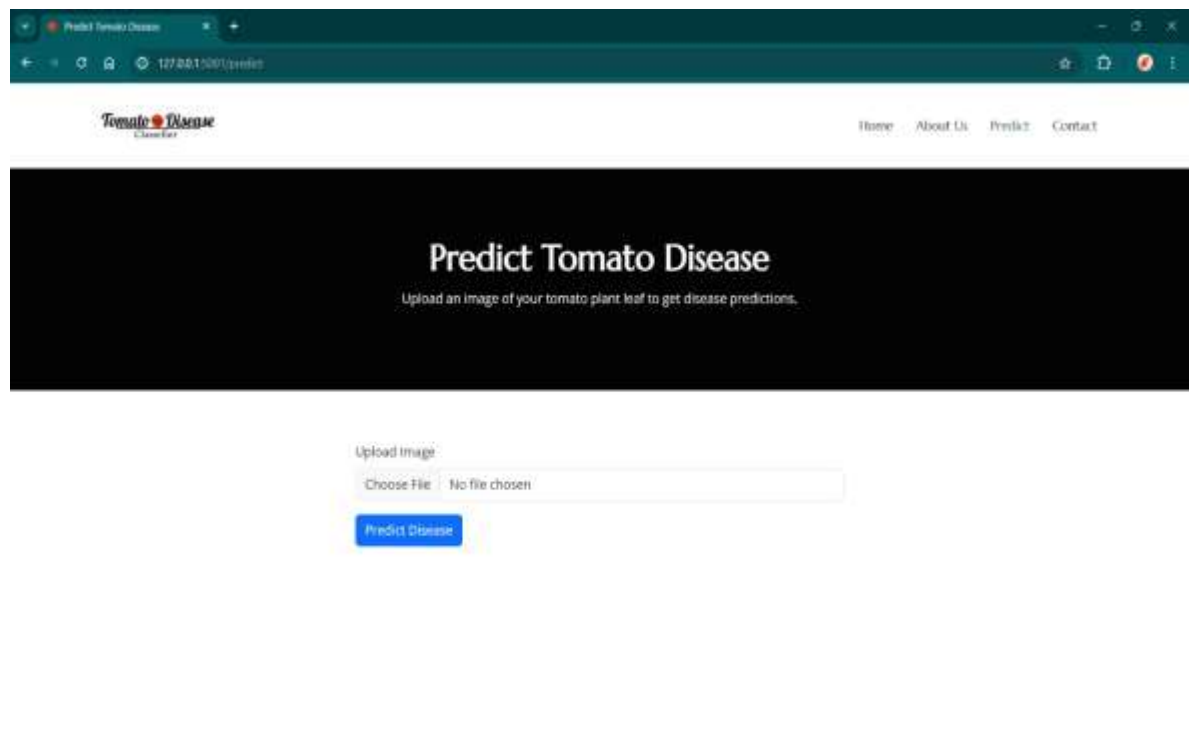
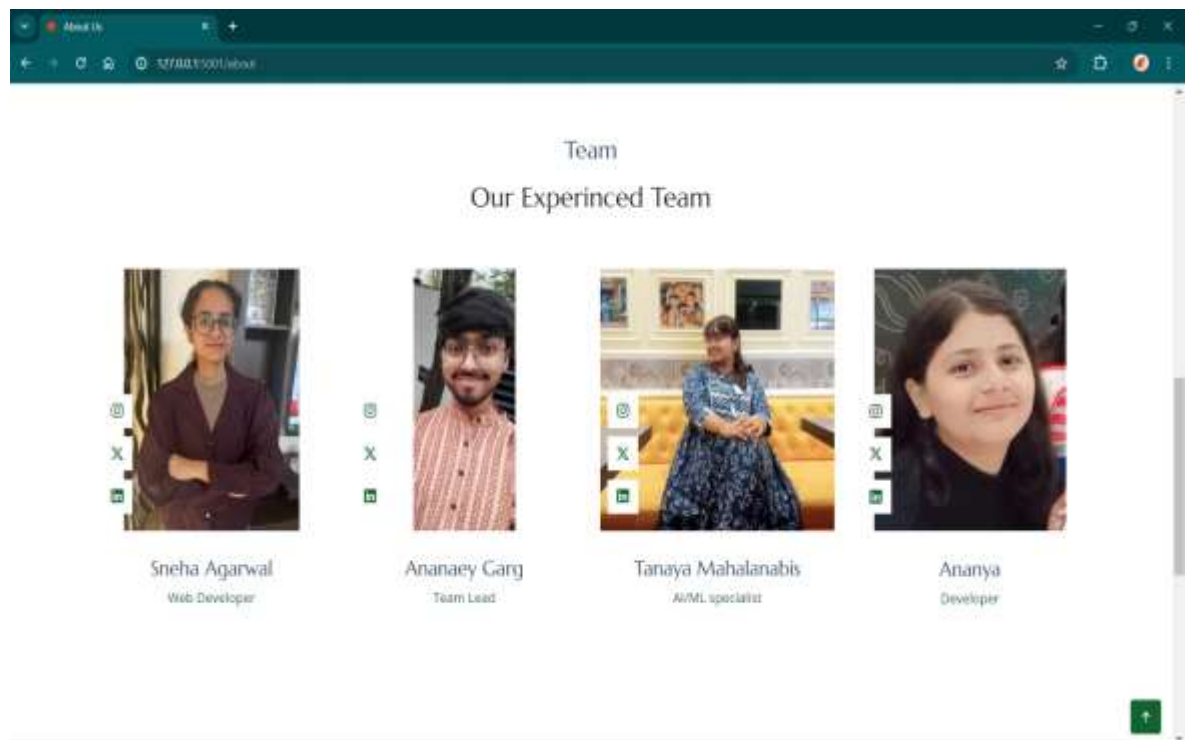
5.2. Final Model Selection Justification

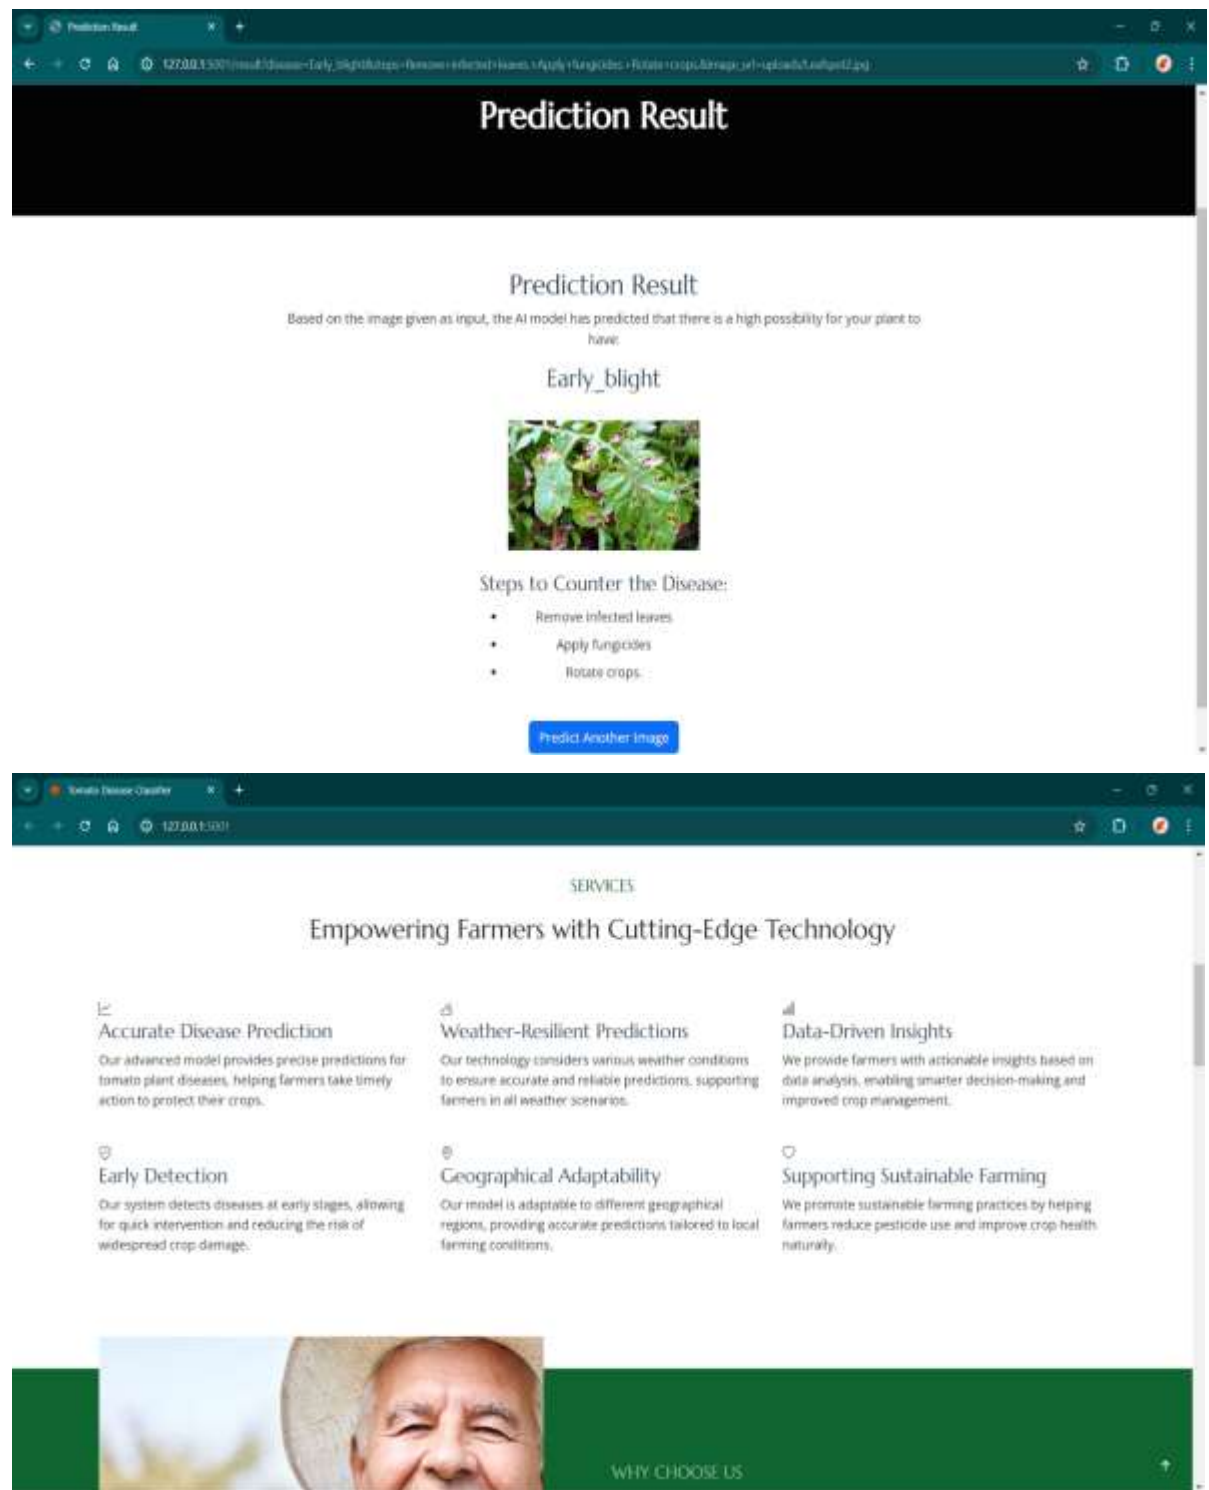
The Final Model Selection Justification centers on the comprehensive evaluation and validation of the ResNet152V2-based architecture for tomato plant disease diagnosis. Following rigorous tuning and training processes, the model consistently demonstrated superior performance in accurately classifying a variety of tomato leaf diseases from images. Its ability to effectively leverage pre-trained features from ImageNet, combined with fine-tuned layers tailored to the specific dataset, ensured robustness and reliability. The decision to select ResNet152V2 was further supported by its scalability, computational efficiency on platforms like Kaggle, and its capability to handle the complexities of plant disease classification tasks. Overall, the model's proven accuracy, efficiency, and adaptability to diverse disease types justified its selection as the final architecture for the diagnostic tool, poised to deliver practical benefits to farmers through enhanced crop management and disease prevention strategies.

6. Results

6.1. Output Screenshots







7. Advantages & Disadvantages

Advantages: The deep learning model based on ResNet152V2 offers significant advantages for tomato plant disease diagnosis. It delivers precise and reliable identification of diseases from leaf images, empowering farmers with timely insights for effective crop management. By reducing reliance on costly laboratory tests, the tool provides a cost-effective solution accessible via smartphones or dedicated devices. This accessibility enhances usability across diverse

farming communities, promoting early intervention and ultimately improving crop yield and economic outcomes. Leveraging platforms like Kaggle ensures scalability and adaptability, positioning the tool as a practical asset for sustainable agriculture.

Disadvantages: Despite its benefits, challenges exist in deploying the ResNet152V2 model. Its performance hinges on the quality and diversity of the training dataset, necessitating careful curation to mitigate biases and ensure accurate disease classification. High computational requirements for model training may pose accessibility barriers in resource-limited settings. Additionally, maintaining model relevance over time requires ongoing updates and integration efforts into existing agricultural practices. Ensuring user trust and acceptance, given the model's complex decision-making process, remains crucial for successful adoption and sustained impact in agricultural communities.

8. Conclusion

In conclusion, the development of a deep learning model based on ResNet152V2 for tomato plant disease diagnosis represents a significant advancement in agricultural technology. This tool promises to revolutionize crop management by providing farmers with accessible, accurate, and timely disease detection capabilities. While challenges such as data quality, computational requirements, and integration complexities exist, the model's potential to enhance crop yield, reduce economic losses, and promote sustainable farming practices outweighs these obstacles. With continued refinement and strategic implementation, this technology holds promise for transforming agriculture, ensuring food security, and fostering economic resilience in farming communities worldwide.

9. Future Scope

Looking ahead, several avenues promise to enhance the impact and capabilities of the ResNet152V2-based tomato plant disease diagnostic tool. Advances in deep learning research could lead to improved model architectures or techniques for handling data biases and increasing interpretability. Integrating real-time disease monitoring using IoT (Internet of Things) devices could provide continuous surveillance and early detection, further optimizing crop management strategies. Collaborations with agronomists and agricultural researchers may expand the dataset to include more diverse environmental conditions and disease variants, enhancing the model's robustness and generalization capabilities. Moreover, leveraging cloud computing and edge AI solutions could address computational challenges, making the tool more accessible to farmers in remote areas. Continuous user feedback and iterative improvements will be essential to tailor the tool to meet evolving agricultural needs, ensuring its relevance and efficacy in sustainable farming practices globally.

10. Appendix

10.1. Source Code


```

# This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input di

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

```

+ Code + Markdown

```

[2]: import os
      os.chdir('/kaggle/input/tomatoleaf/tomato/')
      os.listdir()

```

```

[2]: ['cnn_train.py', 'val', 'train']

```

```

[4]: import warnings
      warnings.filterwarnings('ignore')
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      from PIL import Image
      from tensorflow import keras
      from tensorflow.keras.applications import ResNet152V2
      from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
      from tensorflow.keras.models import Model

```

```

[5]: datagen = keras.preprocessing.image.ImageDataGenerator(rescale=1/255, validation_split=0.3)
      datagen2 = keras.preprocessing.image.ImageDataGenerator(rescale=1/255)

```

```

[6]: #Training and validation dataset
      train = datagen.flow_from_directory('./train', seed=123, subset='training')

      val = datagen.flow_from_directory('./train', seed=123, subset='validation')

      #Test dataset for evaluation
      test = datagen2.flow_from_directory('./val')

```

```

Found 7000 images belonging to 10 classes.
Found 3000 images belonging to 10 classes.
Found 1000 images belonging to 10 classes.

```


[7]:

```
# Training data visualization

classes = os.listdir('./train')

plt.figure(figsize=(25, 10))

for i in enumerate(classes):
    pic = os.listdir('./train/' + i[1])[0]
    image = Image.open('./train/' + i[1] + '/' + pic)
    image = np.asarray(image)
    plt.subplot(2, 5, i[0] + 1)
    plt.title(' {0} / Shape = {1}'.format(i[1], image.shape))
    plt.imshow(image)
plt.show()
```



[8]:

```
base_model = ResNet152V2(input_shape=(256, 256, 3), include_top=False)
for layers in base_model.layers[:148]:
    layers.trainable = False
for layers in base_model.layers[148:]:
    layers.trainable = True
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5

234545216/234545216 — 3s 0us/step

[9]:

```
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1000, activation='relu')(x)
pred = Dense(10, activation='softmax')(x)
model = Model(inputs=base_model.input, outputs=pred)
```

```

#Model Initialization
from tensorflow.keras.applications import ResNet152V2
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
from tensorflow.keras.models import Model

# Define the base model
base_model = ResNet152V2(input_shape=(256,256,3), include_top=False)

# Freeze the layers except the last 140 layers
for layer in base_model.layers[:140]:
    layer.trainable = False

# Unfreeze the last 140 layers
for layer in base_model.layers[140:]:
    layer.trainable = True

# Add custom layers on top of the base model
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1000, activation='relu')(x)
pred = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=pred)

# Print the model summary
model.summary()

```

```
[13]: model.compile(loss='categorical_crossentropy', optimizer='sgd', metrics=['accuracy'])
```

```
[15]: #Model Training
model.fit(train, batch_size= 80, epochs= 15, validation_data=val )
```

```

Epoch 1/15
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR.
10000 00:00:17.21155362.091157 130 device_compiler.h:196] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.
219/219 ----- 193s 633ms/step - accuracy: 0.6414 - loss: 1.1476 - val_accuracy: 0.9280 - val_loss: 0.2212
Epoch 2/15
219/219 ----- 82s 373ms/step - accuracy: 0.9769 - loss: 0.8874 - val_accuracy: 0.9670 - val_loss: 0.1884
Epoch 3/15
219/219 ----- 82s 373ms/step - accuracy: 0.9949 - loss: 0.8243 - val_accuracy: 0.9723 - val_loss: 0.8869
Epoch 4/15
219/219 ----- 82s 373ms/step - accuracy: 0.9974 - loss: 0.8136 - val_accuracy: 0.9717 - val_loss: 0.8872
Epoch 5/15
219/219 ----- 82s 373ms/step - accuracy: 0.9984 - loss: 0.8188 - val_accuracy: 0.9730 - val_loss: 0.8792
Epoch 6/15
219/219 ----- 82s 373ms/step - accuracy: 0.9985 - loss: 0.8076 - val_accuracy: 0.9713 - val_loss: 0.8854
Epoch 7/15
219/219 ----- 82s 373ms/step - accuracy: 0.9988 - loss: 0.8064 - val_accuracy: 0.9767 - val_loss: 0.8678
Epoch 8/15
219/219 ----- 82s 373ms/step - accuracy: 0.9993 - loss: 0.8044 - val_accuracy: 0.9753 - val_loss: 0.8749
Epoch 9/15
219/219 ----- 82s 373ms/step - accuracy: 0.9995 - loss: 0.8037 - val_accuracy: 0.9770 - val_loss: 0.8654
Epoch 10/15
219/219 ----- 82s 373ms/step - accuracy: 0.9999 - loss: 0.8022 - val_accuracy: 0.9793 - val_loss: 0.8645
Epoch 11/15
219/219 ----- 82s 373ms/step - accuracy: 0.9992 - loss: 0.8023 - val_accuracy: 0.9790 - val_loss: 0.8657
Epoch 12/15
219/219 ----- 82s 372ms/step - accuracy: 1.0000 - loss: 0.8016 - val_accuracy: 0.9767 - val_loss: 0.8677
Epoch 13/15
219/219 ----- 82s 373ms/step - accuracy: 0.9998 - loss: 0.8025 - val_accuracy: 0.9803 - val_loss: 0.8677

```

```
[14]: # Prediction and visualizations
classes = os.listdir('./val')
plt.figure(figsize=(18, 28))

for i, cls in enumerate(classes):
    pic_list = os.listdir('./val/' + cls)
    pic = pic_list[np.random.randint(len(pic_list) - 1)]
    image = Image.open('./val/' + cls + '/' + pic)
    image = np.asarray(image)
    pred = np.argmax(model.predict(image.reshape(-1, 256, 256, 3) * 255))

    for j, key in enumerate(list(test.class_indices.keys())):
        if pred == j:
            prediction = key

    plt.subplot(5, 2, i + 1) # Corrected subplot indexing
    plt.title('Actual: {} / Predicted: {}'.format(cls, prediction))
    plt.imshow(image)

plt.show()
```

1/1 ----- 12s 12s/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 10ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 30ms/step
1/1 ----- 0s 10ms/step
1/1 ----- 0s 34ms/step
1/1 ----- 0s 33ms/step
1/1 ----- 0s 33ms/step

Actual: Tomato_Late_blight / Predicted: Tomato_Late_blight

Actual: Tomato_healthy / Predicted: Tomato_healthy

```
model.save('/kaggle/working/tmt.keras')
```

10.2. GitHub & Project Demo Link

GitHub Link: <https://github.com/Frozenfever4/Tomato-Plant-Disease-Detection-From-Leaf-Images>

Demo Drive Link:

<https://drive.google.com/drive/folders/1Wfi51G3vNSOr3rL20mufJvZklKjFs7q4?usp=sharing>